

# INTRODUCTION TO SUPERCOMPUTING

## TMA4280 · Project I

**This project is mandatory, counts for 10% of the final grade and can be done in pairs or alone.**

### Instructions

The deadline is set on the **7th of March 2018**.

The deliverable consists of:

1. a report describing your solutions and interpretation of the results, handed out through Blackboard in PDF format,
2. a GIT repository containing the source code developed to perform the computations.

Practical requirements regarding the code:

- it should be hosted on Github,
- it should be structured in different subdirectories addressing the different questions,
- it should contain results gathered in text files,
- it must be written in C, C++ using C arrays only, or FORTRAN,
- it must use double precision,
- it must compile and run using Makefile targets,
- results presented in the report have to be reproducible.

As soon as you have decided whether you want to work in pairs or alone:

- create a Github repository (one per pair) named **TMA4280v2018**,
- send your name(s) and the link to the repository by email.

All the developments for this project will be contained in a subdirectory named **P1**.

## Description

The project deals with the computation of  $\pi$  using different approximations and their parallelization using first MPI, and then with a small modification to use OpenMP.

Considered approximations are based on the development in series,

$$S_n = \sum_{i=1}^n v_i.$$

with  $v_i$  depending on the chosen method.

**1. Riemann zeta function**  $\zeta(s)$  with  $s = 2$ : Each term is defined as:

$$v_i = \frac{1}{i^2}, \quad i = 1, \dots, n$$

and

$$S = \lim_{n \rightarrow \infty} S_n = \frac{\pi^2}{6}.$$

**2. Machin Formula:** Given  $x \in [-1, 1]$ ,

$$v_i(x) = (-1)^{i-1} \frac{x^{2i-1}}{2i-1}$$

and

$$S(x) = \lim_{n \rightarrow \infty} S_n(x) = \arctan(x)$$

with

$$\frac{\pi}{4} = 4 \arctan\left(\frac{1}{5}\right) - \arctan\left(\frac{1}{239}\right)$$

In all the questions, the name in bold and enclosed in parenthesis denotes the subdirectory where the implementation should be located. All the computations should run with the command `make test` and in the parallel case specifying the number of processes with the environment variable `NP`. You may factor common code in a library if you want, but it is not a requirement.

*Question 1. Serial implementation.* Write a serial program implementing the computation of  $\pi$  for a given  $n$  read from the command line for:

- Method 1. add a program in (**zeta0**)
- Method 2. add a program in (**mach0**)

*Question 2. Unit test.* Every development should come with unit testing to check the logic of the implementation. Such tests should execute quickly and compare a computed value against an expected value. Implement a small test comparing the value of each series with  $n = 3$ .

- Method 1. add a unit test in (**zeta0**)
- Method 2. add a unit test in (**mach0**)

The test should be implemented in a simple function (no unit test framework required) and executed with `make utablest`. Why do you think such test may be useful when parallelizing a computational code?

*Question 3. Verification test.* When the logic of the implementation is tested then the mathematical properties of the algorithm should be assessed: convergence, stability. Such tests should execute fairly quickly and verify the behaviour of the algorithm as compared to the theory if possible. Implement a small test computing the error  $|\pi - \pi_n|$  for  $n = 2^k$  with  $k = 1, \dots, 24$ :

- Method 1. add a verification test in (**zeta0**)
- Method 2. add a verification test in (**mach0**)

The test should be executed with `make vtest` and results should be saved in a file. Comment on the obtained results.

*Question 4. Data distribution.* As we are interested in computing the sum of all vector elements  $v_i$  numerically, we will work under the constraint that the values should be put in a vector before being summed. The suggest program deliberately relies on partitioning and distribution of the data by process zero.

Process zero **only** should be responsible for:

- generating the vector elements,
- partitioning the vector in a way that the problem is load-balanced,
- and distribute the elements to all the processes.

Each process will work on the received data.

- Method 1. add a program in (**zeta1**)
- Method 2. add a program in (**mach1**)

Can you comment on the limitation of such approach for the data distribution and a possible improvement? Provide arguments to support your answer.

*Question 5. MPI implementation.* Modify further the program to compute the approximation of  $\pi$  using both methods such that each process:

- computes a partial sum from its data,
- then all the partial sums should be added together on the root process,
- and then global sum is printed on the standard output by the root process.

Only the root process holds the final value. Report the error  $|\pi - \pi_n|$  in double precision for different values of  $n$  and the wall time, for different number of MPI processes which are powers of two. The program should contain an assertion and fail if the number of processes is not a power of two, as a design constraint.

- Method 1. modify the program and add a test in (**zeta1**)
- Method 2. modify the program and add a test in (**mach1**)

Plot the error and the timings. Which MPI calls were convenient and/or necessary to use? Can you comment on the methodology used for computing the wall time?

*Question 6. Analysis*

Compare the errors from the single-process program and the multi-process program for  $P = 2$  and  $P = 8$ . Should the answer be the same in all cases? Exactly, or approximately? Can you explain why? Provide arguments to support your answer.

*Question 7. Global reduction.* Modify the final step, the reduction with `MPI_SUM`, such that all processes store the global sum: first by using an MPI function and then by implementing the recursive-doubling sum.

- Add a program in (**reduc**)

Do a small scaling study, what do you observe?

*Question 8. OpenMP implementation.* Make the necessary changes needed to use shared memory parallelization with OpenMP.

- Method 1. add a program in (**zeta2**)
- Method 2. add a program in (**mach2**)

Perform the same analysis as for the *MPI implementation*.

*Question 9. Hybrid MPI/OpenMP implementation.* Confirm that your program also works when using OpenMP and MPI in combination. What is the advantage of running such configuration?

*Question 10. Discussion.*

- Compare the memory requirement per process for the single-process program and the multi-process program when  $n \gg 1$ .
- How many floating point operations are needed to generate a vector  $v$ ?
- How many are needed to compute  $S_n$ ?
- Is the multi-process program load balanced?

*Question 11. Conclusion.* Do you consider parallel processing attractive for solving this problem? Explain why.